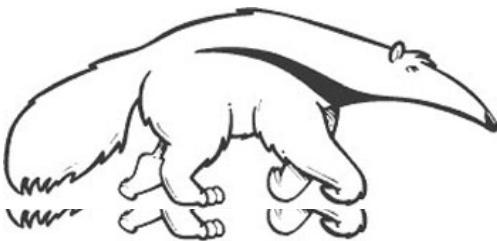


+

Machine Learning and Data Mining

Introduction

Prof. Alexander Ihler



Artificial Intelligence (AI)

- Building “intelligent systems”
- Lots of parts to intelligent behavior



RoboCup





Chess (Deep Blue v. Kasparov)



Darpa GC (Stanley)

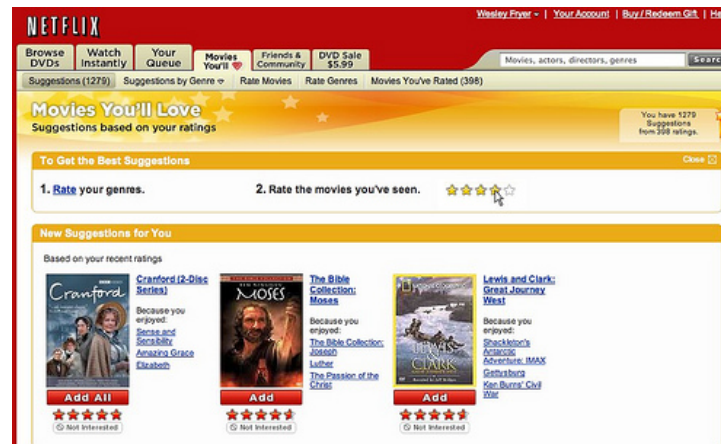

Machine learning (ML)

- One (important) part of AI
- Making predictions (or decisions)
- Getting better with experience (data)
- Problems whose solutions are “hard to describe”





Types of prediction problems

- Supervised learning
 - “Labeled” training data
 - Every example has a desired target value (a “best answer”)
 - Reward prediction being close to target
- Classification: a discrete-valued prediction (often: action / decision)
- Regression: a continuous-valued prediction



Types of prediction problems

- Supervised learning
- Unsupervised learning
 - No known target values
 - No targets = nothing to predict?
 - Reward “patterns” or “explaining features”
 - Often, data mining




Types of prediction problems

- Supervised learning
- Unsupervised learning
- Semi-supervised learning
 - Similar to supervised
 - some data have unknown target values
- Ex: medical data
 - Lots of patient data, few known outcomes
- Ex: image tagging
 - Lots of images on Flickr, but only some of them tagged

Types of prediction problems

- Supervised learning
- Unsupervised learning
- Semi-supervised learning
- Reinforcement learning
- “Indirect” feedback on quality
 - No answers, just “better” or “worse”
 - Feedback may be delayed



Logistics

- Course webpage for assignments & other info (?)
- EEE for homework submission & return
- Piazza for questions & discussions
- No required textbook
 - Recommended: Murphy, “Machine Learning...”, 2012.
 - Also
 - Duda, Hart & Stork, “Pattern classification”
 - Hastie, Tibshirani & Friedman, “Elements of Statistical Learning”
- But
 - I’ll try to cover everything needed in lectures and notes
 - All textbooks mainly for reference purposes

Logistics

- Grading (may be subject to change)
 - 25% homework (6+? >5: drop 1)
 - 15% project (Kaggle)
 - 25% midterm, 35% final
 - Due 11:59pm listed day, EEE or my office
 - Late homework:
 - 10% off per day or part
 - No credit after solutions posted: turn in what you have
- Collaboration
 - Study groups, discussion, assistance encouraged
 - Whiteboards, etc.
 - Do your homework yourself
 - Don't exchange solutions or HW code

Data exploration

- Machine learning is a data science
 - Look at the data; get a “feel” for what might work
- What types of data do we have?
 - Binary values? (spam; gender; ...)
 - Categories? (home state; labels; ...)
 - Integer values? (1..5 stars; age brackets; ...)
 - (nearly) real values? (pixel intensity; prices; ...)
- Are there missing data?
- “Shape” of the data? Outliers?

Scientific software

- **Python**
 - Numpy, Matplotlib, SciPy...
- Matlab
 - Octave (free)
- R
 - Used mainly in statistics
- C++
 - For performance, not prototyping
- And other, more specialized languages for modeling...

Representing data

- Example: Fisher's "Iris" data
http://en.wikipedia.org/wiki/Iris_flower_data_set
- Three different types of iris
 - "Class", y
- Four "features", x_1, \dots, x_4
 - Length & width of sepals & petals
- 150 examples (data points)



Representing the data (Matlab)

- Have m observations (data points)

$$\{x^{(1)} \dots, x^{(m)}\}$$

- Each observation is a vector consisting of n features

$$x^{(j)} = [x_1^{(j)} x_2^{(j)} \dots x_n^{(j)}]$$

- Often, represent this as a “data matrix”

$$\underline{X} = \begin{bmatrix} x_1^{(1)} & \dots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_1^{(m)} & \dots & x_n^{(m)} \end{bmatrix}$$

```
import numpy as np    # import numpy
iris = np.genfromtxt("data/iris.txt", delimiter=None)
X = iris[:,0:4]        # load data and split into features, targets
Y = iris[:,4]
print X.shape          # 150 data points; 4 features each
(150, 4)
```


Basic statistics

- Look at basic information about features
 - Average value? (mean, median, etc.)
 - “Spread”? (standard deviation, etc.)
 - Maximum / Minimum values?

```
print np.mean(X, axis=0)      # compute mean of each feature
[ 5.8433  3.0573  3.7580  1.1993 ]
print np.std(X, axis=0)      #compute standard deviation of each feature
[ 0.8281  0.4359  1.7653  0.7622 ]
print np.max(X, axis=0)      # largest value per feature
[ 7.9411  4.3632  6.8606  2.5236 ]
print np.min(X, axis=0)      # smallest value per feature
[ 4.2985  1.9708  1.0331  0.0536 ]
```

Histograms

- Count the data falling in each of K bins
 - “Summarize” data as a length-K vector of counts (& plot)
 - Value of K determines “summarization”; depends on # of data
 - K too big: every data point falls in its own bin; just “memorizes”
 - K too small: all data in one or two bins; oversimplifies



% Histograms in Matplotlib

```
import matplotlib.pyplot as plt
```

```
X1 = X[:,0]
```

```
Bins = np.linspace(4,8,17)
```

```
plt.hist( X1, bins=Bins )
```


```
# extract first feature
```

```
# use explicit bin locations
```

```
# generate the plot
```


Scatterplots

- Illustrate the relationship between two features

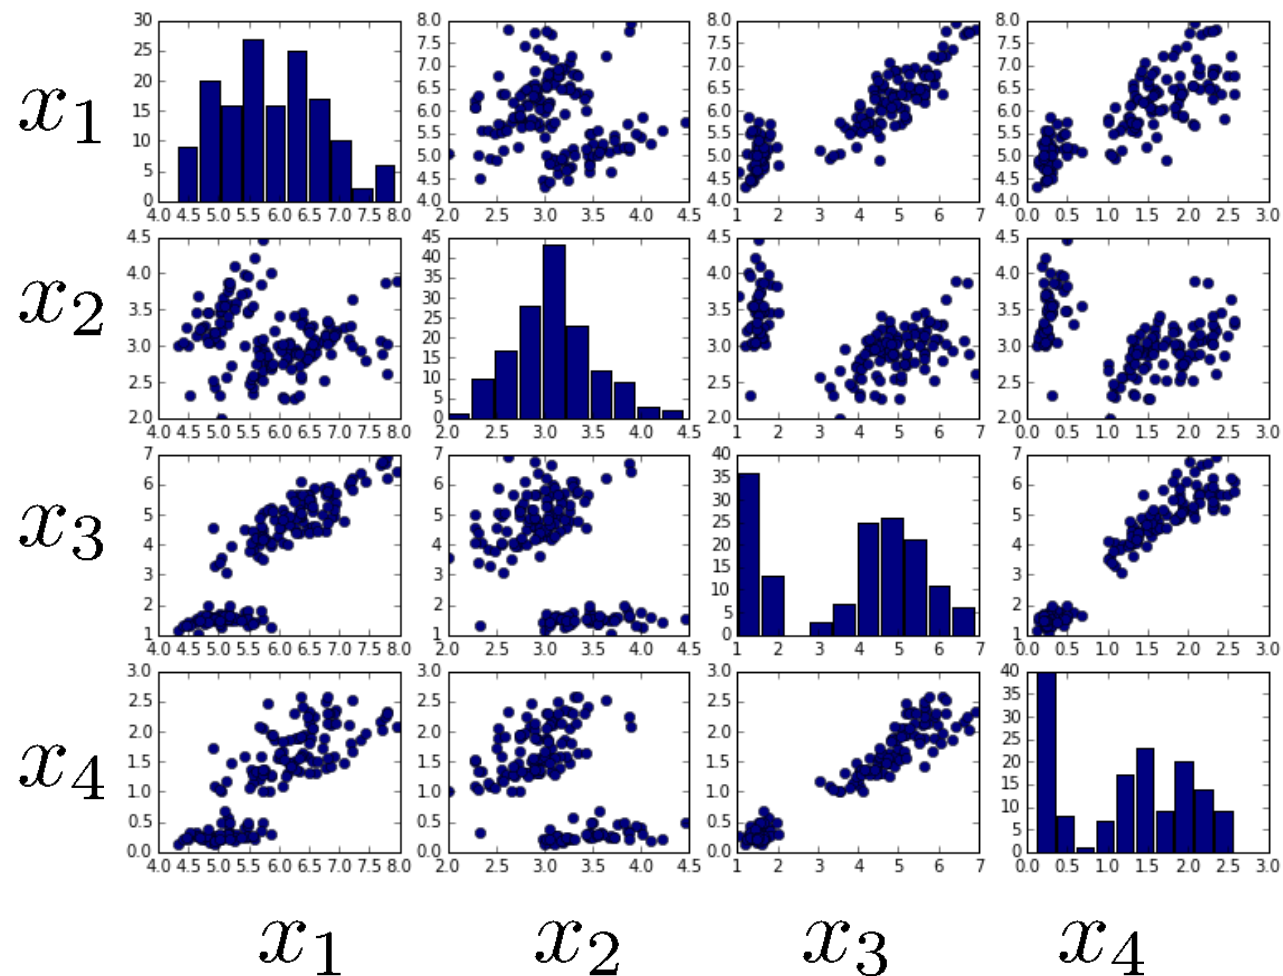


```
% Plotting in Matplotlib  
plt.plot(X[:,0], X[:,1], 'b.');
```

% plot data points as blue dots


Scatterplots

- For more than two features we can use a pair plot:




Supervised learning and targets


- Supervised learning: predict target values
- For discrete targets, often visualize with color



```
plt.hist( [X[Y==c,1] for c in np.unique(Y)] ,  
          bins=20, histtype='barstacked')
```



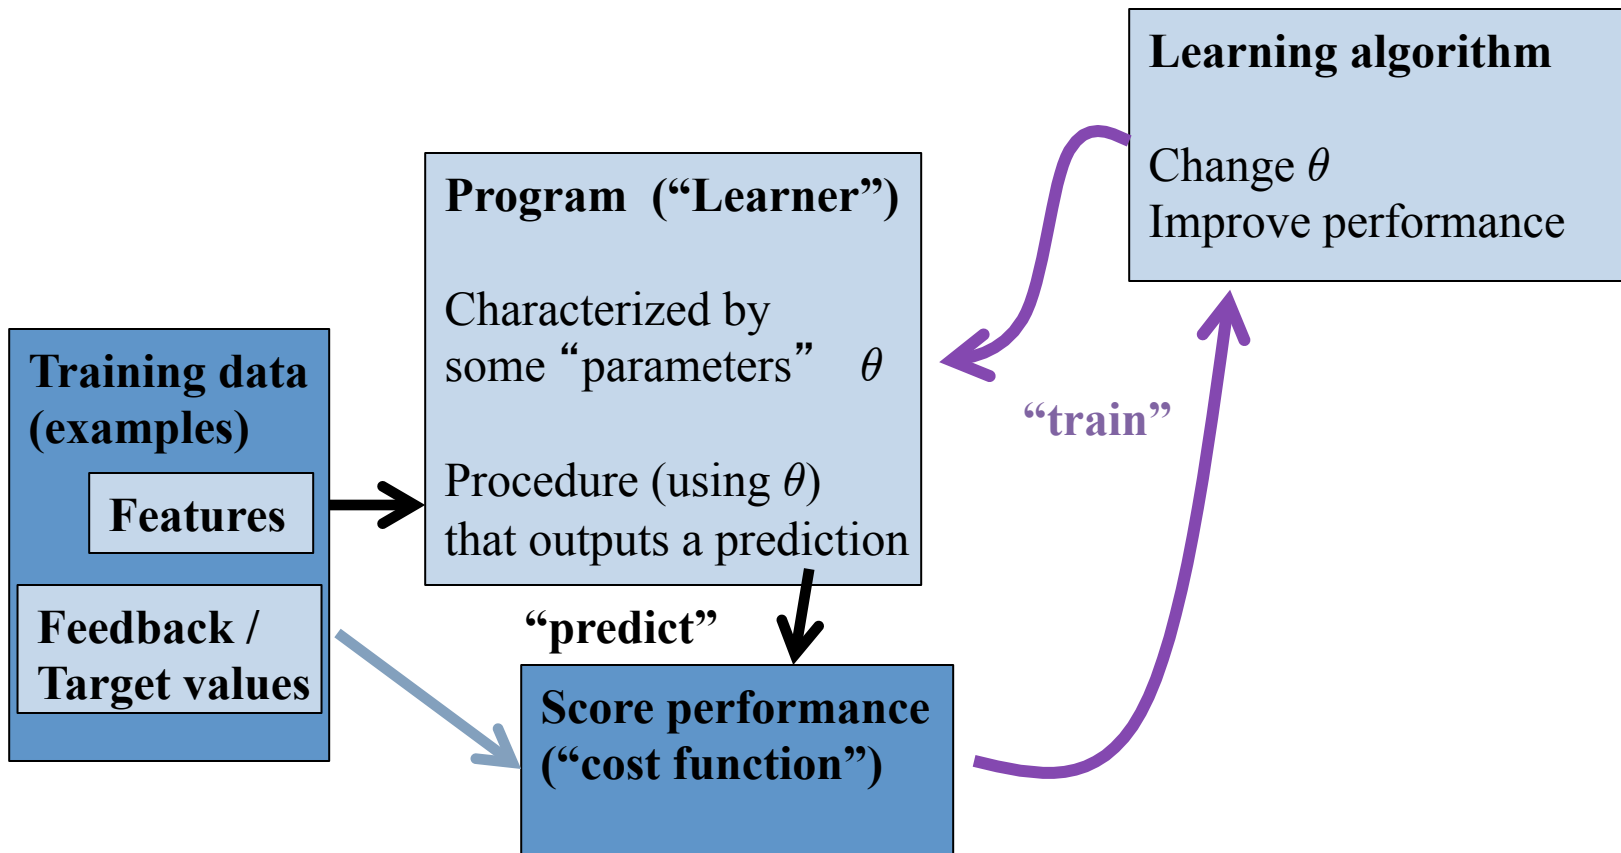
```
ml.histy(X[:,1], Y, bins=20)
```



```
colors = ['b','g','r']  
for c in np.unique(Y):  
    plt.plot( X[Y==c,0], X[Y==c,1], 'o',  
              color=colors[int(c)] )
```

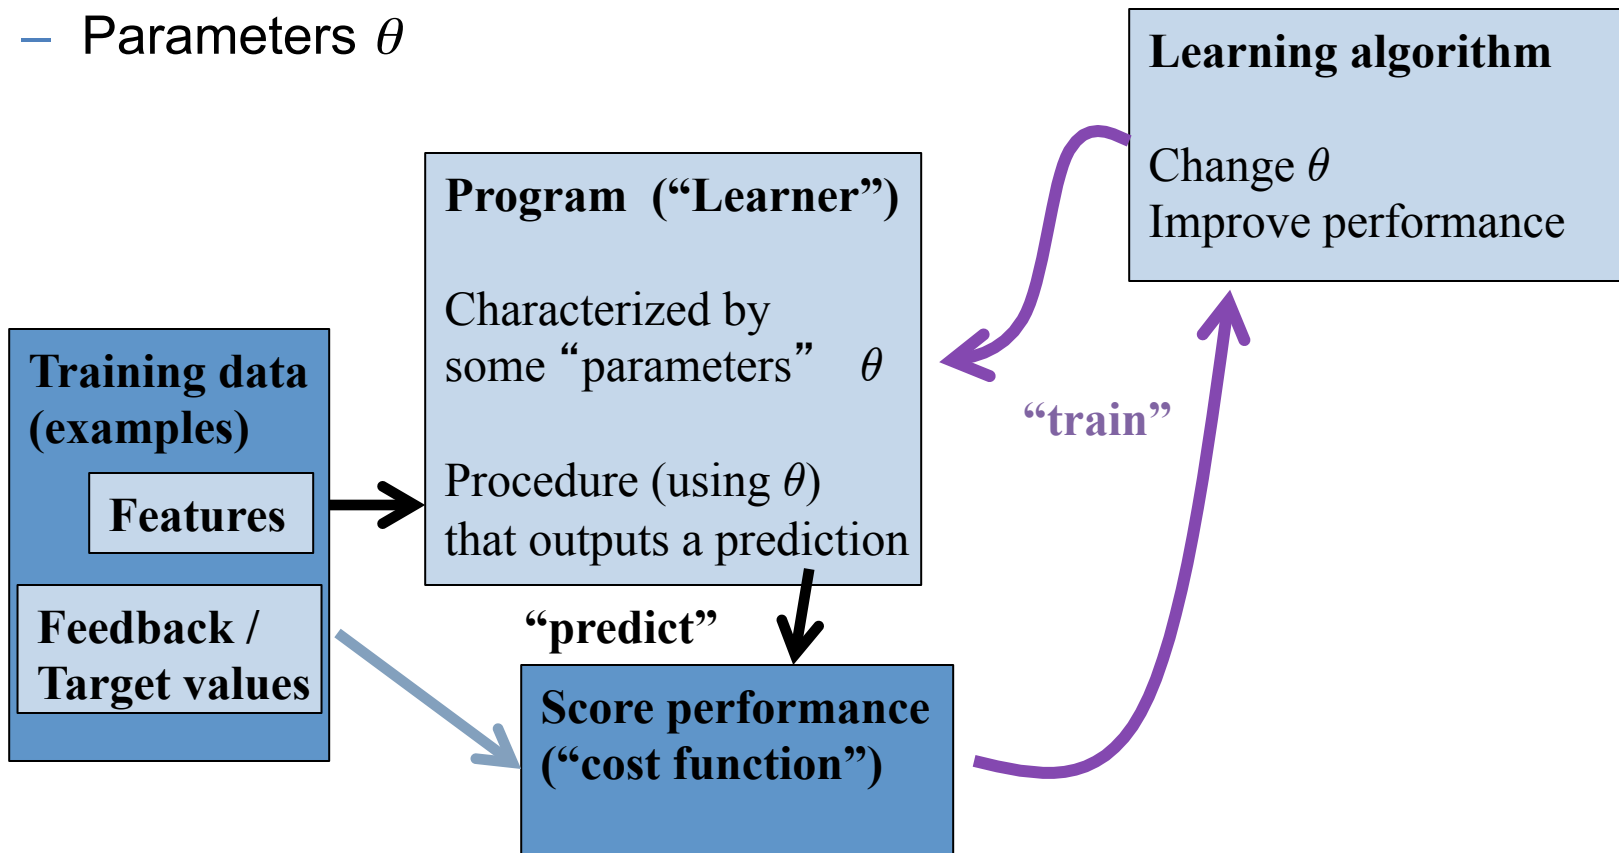
How does machine learning work?

- “Meta-programming”
 - Predict – apply rules to examples
 - Score – get feedback on performance
 - Learn – change predictor to do better




Supervised learning

- Notation
 - Features x
 - Targets y
 - Predictions $\hat{y} = f(x ; \theta)$
 - Parameters θ




Regression; Scatter plots




- Suggests a relationship between x and y
- *Prediction*: new x, what is y?

Nearest neighbor regression



- Find training datum $x^{(i)}$ closest to $x^{(new)}$
Predict $y^{(i)}$

Nearest neighbor regression




“Predictor”:

Given new features:
Find nearest example
Return its value

- Defines a function $f(x)$ implicitly
- “Form” is piecewise constant

Linear regression



“Predictor”:


Evaluate line:

$$r = \theta_0 + \theta_1 x_1$$

return r

- Define form of function $f(x)$ explicitly
- Find a good $f(x)$ within that family


Measuring error



$$\text{MSE} = \frac{1}{m} \sum_i (y^{(i)} - \hat{y}(x^{(i)}))^2$$

Regression vs. Classification

Regression




Features x

Real-valued target y

Predict continuous function $\hat{y}(x)$

Classification




Features x

Discrete class c


(usually 0/1 or +1/-1)

Predict discrete function $\hat{y}(x)$


Classification



Classification



Measuring error



$$\text{ERR} = \frac{1}{m} \sum_i [y^{(i)} \neq \hat{y}(x^{(i)})]$$

A simple, optimal classifier

- Classifier $f(x; \theta)$
 - maps observations x to predicted target values
- Simple example
 - Discrete feature x : $f(x; \theta)$ is a contingency table
 - Ex: spam filtering: observe just $X_1 =$ in contact list?
- Suppose we knew the true conditional probabilities:
- Best prediction is the most likely target!

Feature	spam	keep
X=0	0.6	0.4
X=1	0.1	0.9

“Bayes error rate”

$$\begin{aligned} & \Pr[X=0] * \Pr[\text{wrong} \mid X=0] + \Pr[X=1] * \Pr[\text{wrong} \mid X=1] \\ &= \Pr[X=0] * (1 - \Pr[Y=S \mid X=0]) + \Pr[X=1] * (1 - \Pr[Y=K \mid X=1]) \end{aligned}$$

Optimal least-squares regression

- Suppose that we know true $p(X,Y)$
- Prediction $f(x)$: *arbitrary* function
 - Focus on some specific x : $f(x) = v$

- Expected squared error loss is

$$\mathbb{E}_{Y|X=x}[(Y - v)^2] = \int p(Y|X = x)(Y - v)^2 dY$$

- Minimum: take derivative & set to zero

$$\frac{\partial}{\partial v} \int p(Y|X = x)(Y - v)^2 dY = \int p(Y|X = x)2(Y - v) = 0$$


$$\Rightarrow 2 \int p(Y|X = x)Y = 2 \left(\int p(Y|X = x) \right) v$$

$$\Rightarrow v = \int p(Y|X = x)Y = \mathbb{E}_{Y|X=x}[Y]$$

Optimal estimate of Y: conditional expectation given X

Bayes classifier, estimated


- Now, let's see what happens with “real” data
 - Use empirically estimated probability **model** for $p(x,y)$
- Iris data set, first feature only (real-valued)
 - We can estimate the probabilities (e.g., with a histogram)



2 Bins:


Predict “green” if $X < 3.25$, else “blue”

Model is “too simple”



20 Bins:

Predict by majority color in each bin



500 Bins:



Each bin has ~ 1 data point!

What about bins with 0 data?


Model is “too complex”

Inductive bias


- “Extend” observed data to unobserved examples
 - “Interpolate” / “extrapolate”
- What kinds of functions to expect? Prefer these (“bias”)
 - Usually, let data pull us away from assumptions only with evidence!




Overfitting and complexity




Overfitting and complexity




Overfitting and complexity




Overfitting and complexity



Overfitting and complexity



How Overfitting affects Prediction



Competitions

- Training data
 - Used to build your model(s)
- Validation data
 - Used to assess, select among, or combine models
 - Personal validation; leaderboard; ...
- Test data
 - Used to estimate “real world” performance

#	Δ1w	Team Name <small>* in the money</small>	Score <small>?</small>	Entries	Last Submission U1
1	-	BrickMover <small>1*</small>	1.21251	40	Sat, 31 Aug 2013 23:...
2	new	vsu <small>*</small>	1.21552	13	Sat, 31 Aug 2013 20:...
3	↑2	Merlion	1.22724	29	Sat, 31 Aug 2013 23:...
4	↓2	Sergey	1.22856	15	Sat, 31 Aug 2013 23:...
5	new	liuyongqi	1.22980	13	Sat, 31 Aug 2013 13:...

Summary

- What is machine learning?
 - Types of machine learning
 - How machine learning works
- Supervised learning
 - Training data: features x , targets y
- Regression
 - (x,y) scatterplots; predictor outputs $f(x)$; optimal MSE predictor
- Classification
 - (x,x) scatterplots
 - Decision boundaries, colors & symbols; Bayes optimal classifier
- Complexity
 - Training vs test error
 - Under- & over-fitting